

**Recommender systems** or **recommendation systems** (sometimes replacing "system" with a synonym such as platform or engine) are a subclass of [information filtering system](#) that seek to predict the 'rating' or 'preference' that user would give to an item (such as [music](#), [books](#), or [movies](#)) or social element (e.g. [people](#) or [groups](#)) they had not yet considered, using a model built from the characteristics of an item (content-based approaches) or the user's social environment (collaborative filtering approaches).<sup>[1][2]</sup>

Recommender systems have become extremely common in recent years. A few examples of such systems:

- When viewing a product on [Amazon.com](#), the store will recommend additional items based on a matrix of what other shoppers bought along with the currently selected item.<sup>[3]</sup>
- [Pandora Radio](#) takes an initial input of a song or musician and plays music with similar characteristics (based on a series of keywords attributed to the inputted artist or piece of music). The stations created by Pandora can be refined through user feedback (emphasizing or deemphasizing certain characteristics).
- [Netflix](#) offers predictions of movies that a user might like to watch based on the user's previous ratings and watching habits (as compared to the behavior of other users), also taking into account the characteristics (such as the genre) of the film.

## Overview

---

Recommender systems typically produce a list of recommendations in one of two ways - through collaborative or content-based filtering. [Collaborative filtering](#) approaches to build a model from a user's past behavior (items previously purchased or selected and/or numerical ratings given to those items) as well as similar decisions made by other users, then use that model to predict items (or ratings for items) that the user may have an interest in.<sup>[4]</sup> [Content-based filtering](#) approaches utilize a series of discrete characteristics of an item in order to recommend additional items with similar properties.<sup>[5]</sup> These approaches are often combined (see Hybrid Recommender Systems).

The differences between collaborative and content-based filtering can be demonstrated by comparing two popular music recommender systems - [Last.fm](#) and [Pandora Radio](#).

- Pandora uses the properties of a song or artist (a subset of the 400 attributes provided by the [Music Genome Project](#)) in order to seed a "station" that plays music with similar properties. User feedback is used to refine the station's results, deemphasizing certain attributes when a user "dislikes" a particular song and emphasizing other attributes when a user "likes" a song. This is an example of a content-based approach.
- Last.fm creates a "station" of recommended songs by observing what bands and individual tracks that the user has listened to on a regular basis and comparing those against the listening behavior of other users. Last.fm will play tracks that do not appear in the user's library, but are often played by other users with similar interests. As this approach leverages the behavior of users, it is an example of a collaborative filtering technique.

Each type of system has its own strengths and weaknesses. In the above example, Last.fm requires a large amount of information on a user in order to make accurate recommendations. This is an example of the [cold start](#) problem, and is common in collaborative filtering systems.<sup>[6]</sup> While Pandora needs very little information to get started, it is far more limited in scope (for example, it can only make recommendations that are similar to the original seed).

Recommender systems are a useful alternative to [search algorithms](#) since they help users discover items they might not have found by themselves. Interestingly enough, recommender systems are often implemented using search engines indexing non-traditional data.

Montaner provides the first overview of recommender systems, from an intelligent agents perspective.<sup>[7]</sup> Adomavicius provides a new overview of recommender systems.<sup>[8]</sup> Herlocker provides an additional overview of evaluation techniques for recommender systems.<sup>[9]</sup>

Recommender system is an active research area in the [data mining](#) and [machine learning](#) areas. Some conferences such as RecSys, SIGIR, KDD have it as a topic.

## Approaches

---

### [\[edit\]](#) Collaborative filtering

Main article: [Collaborative filtering](#)

One approach to the design of recommender systems that has seen wide use is [collaborative filtering](#).<sup>[10]</sup> Collaborative filtering methods are based on collecting and analyzing a large amount of information on users' behaviors, activities or preferences and predicting what users will like based on their similarity to other users. A key advantage of the collaborative filtering approach is that it does not rely on machine analyzable content and therefore it is capable of accurately recommending complex items such as movies without requiring an "understanding" of the item itself. Many algorithms have been used in measuring user similarity or item similarity in recommender systems. For example, the [k-nearest neighborhood](#) (k-NN) approach<sup>[11]</sup> and the [Pearson Correlation](#).

When building a model from a user's profile, a distinction is often made between explicit and [implicit](#) forms of [data collection](#).

Examples of [explicit data collection](#) include the following:

- Asking a user to rate an item on a sliding scale.
- Asking a user to rank a collection of items from favorite to least favorite.
- Presenting two items to a user and asking him/her to choose the better one of them.
- Asking a user to create a list of items that he/she likes.

Examples of [implicit data collection](#) include the following:

- Observing the items that a user views in an online store.
- Analyzing item/user viewing times<sup>[12]</sup>
- Keeping a record of the items that a user purchases online.
- Obtaining a list of items that a user has listened to or watched on his/her computer.
- Analyzing the user's social network and discovering similar likes and dislikes

The recommender system compares the collected data to similar and dissimilar data collected from others and calculates a list of recommended items for the user. Several commercial and non-commercial examples are listed in the article on [collaborative filtering systems](#).

One of the most famous examples of collaborative filtering is item-to-item collaborative filtering (people who buy x also buy y), an algorithm popularized by [Amazon.com](#)'s recommender system.<sup>[3]</sup> Other examples include:

- As previously detailed, [Last.fm](#) recommends music based on a comparison of the listening habits of similar users.
- [Facebook](#), [MySpace](#), [LinkedIn](#), and other social networks use collaborative filtering to recommend new friends, groups, and other social connections (by examining the network of connections between a user and their friends).<sup>[4]</sup>

Collaborative filtering approaches often suffer from three problems: [cold start](#), scalability, and sparsity.<sup>[13]</sup>

- Cold Start: These systems often require a large amount of existing data on a user in order to make accurate recommendations.
- Scalability: In many of the environments that these systems make recommendations in, there are millions of users and products. Thus, a large amount of computation power is often necessary to calculate recommendations.
- Sparsity: The number of items sold on major e-commerce sites is extremely large. The most active users will only have rated a small subset of the overall database. Thus, even the most popular items have very few ratings.

A particular type of collaborative filtering algorithm uses [matrix factorization](#), a [low-rank matrix approximation](#) technique.<sup>[14][15][16]</sup>

## [\[edit\]](#) **Content-based filtering**

*[citation needed]*

Another common approach when designing recommender systems is [content-based filtering](#). Content-based filtering methods are based on information about and characteristics of the items that are going to be recommended. In other words, these algorithms try to recommend items that are similar to those that a user liked in the past (or is examining in the present). In particular, various candidate items are compared with items previously rated by the user and the best-matching items are recommended. This approach has its roots in [information retrieval](#) and [information filtering](#) research.

Basically, these methods use an item profile (i.e., a set of discrete attributes and features) characterizing the item within the system. The system creates a content-based profile of users based on a weighted vector of item features. The weights denote the importance of each feature to the user and can be computed from individually rated content vectors using a variety of

techniques. Simple approaches use the average values of the rated item vector while other sophisticated methods use machine learning techniques such as [Bayesian Classifiers](#), [cluster analysis](#), [decision trees](#), and [artificial neural networks](#) in order to estimate the probability that the user is going to like the item.

Direct feedback from a user, usually in the form of a [like or dislike button](#), can be used to assign higher or lower weights on the importance of certain attributes (using [Rocchio Classification](#) or other similar techniques).

A key issue with content-based filtering is whether the system is able to learn user preferences from user's actions regarding one content source and use them across other content types. When the system is limited to recommending content of the same type as the user is already using, the value from the recommendation system is significantly less than when other content types from other services can be recommended. For example, recommending news articles based on browsing of news is useful, but it's much more useful when music, videos, products, discussions etc. from different services can be recommended based on news browsing. Such cross-content recommendation has been productised by [Leiki](#).

As previously detailed, Pandora Radio is a popular example of a content-based recommender system that plays music with similar characteristics to that of a song provided by the user as an initial seed. There are also a large number of content-based recommender systems aimed at providing movie recommendations, a few such examples include [Rotten Tomatoes](#), [Internet Movie Database](#), [Jinni](#), [Rovi Corporation](#) and [See This Next](#).

### [\[edit\]](#) **Hybrid Recommender Systems**

Recent research has demonstrated that a hybrid approach, combining [collaborative filtering](#) and [content-based filtering](#) could be more effective in some cases. Hybrid approaches can be implemented in several ways: by making content-based and collaborative-based predictions separately and then combining them; by adding content-based capabilities to a collaborative-based approach (and vice versa); or by unifying the approaches into one model (see [\[8\]](#) for a complete review of recommender systems). Several studies empirically compare the performance of the hybrid with the pure collaborative and content-based methods and demonstrate that the hybrid methods can provide more accurate recommendations than pure approaches. These methods can also be used to overcome some of the common problems in recommender systems such as cold start and the sparsity problem.

[Netflix](#) and [See This Next](#) are good examples of hybrid systems. They make recommendations by comparing the watching and searching habits of similar users (i.e. collaborative filtering) as well as by offering movies that share characteristics with films that a user has rated highly (content-based filtering).

[\[edit\]](#)